


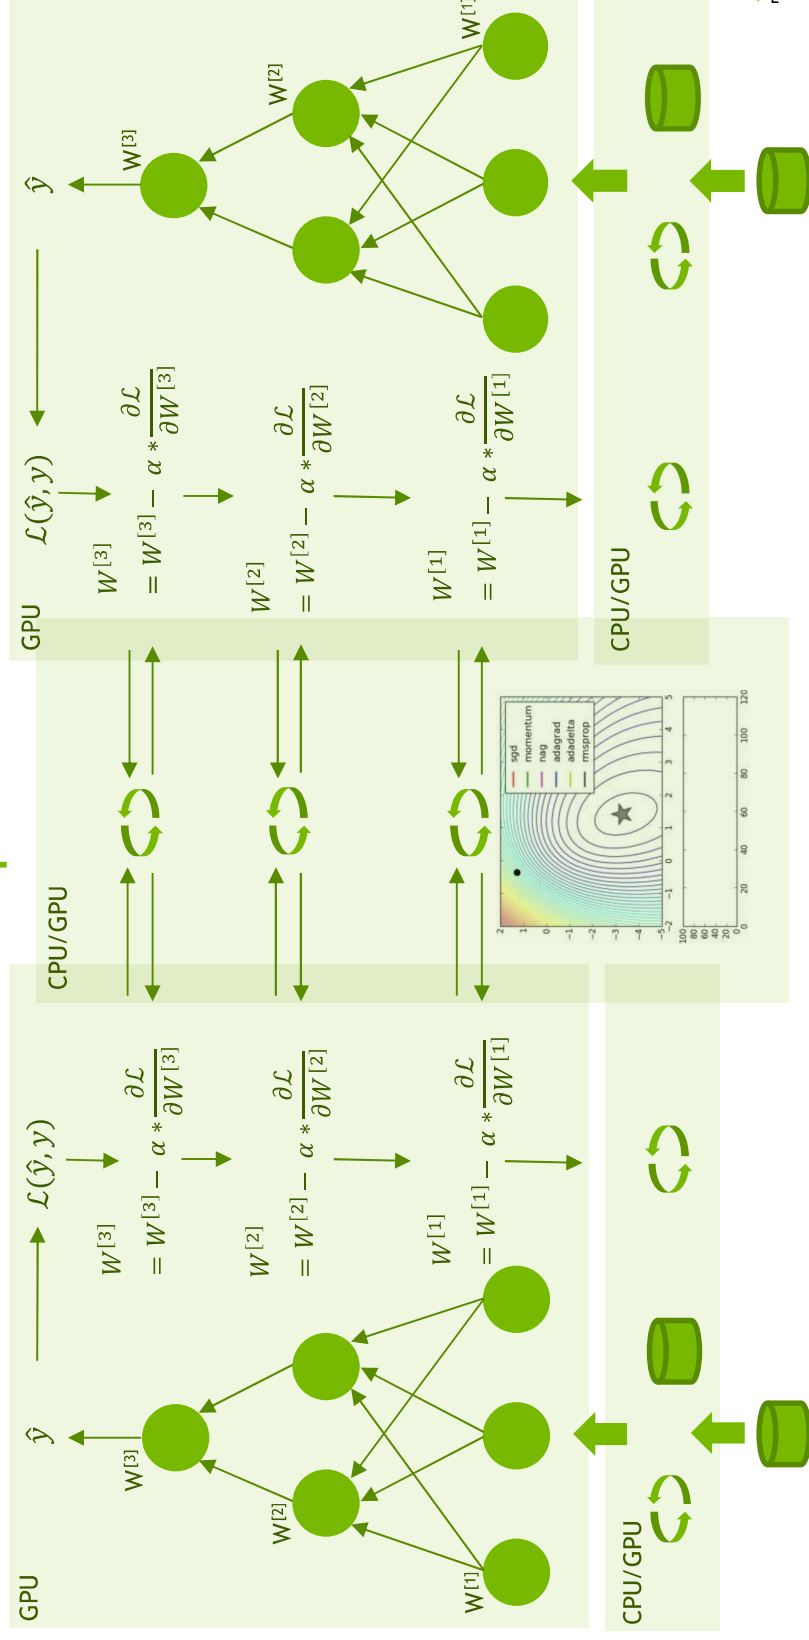
# DATA PARALLELISM: HOW TO TRAIN DEEP LEARNING MODELS ON MULTIPLE GPUS

LAB 2, PART 1: INTRODUCTION TO DISTRIBUTED DATA PARALLEL (DDP)



# TRAINING A NEURAL NETWORK

## Multiple GPUs



# MEET DDP

Library for distributed DL

Prepackaged into and optimized for  
PyTorch, an increasingly popular platform  
among ML engineers and researchers



# USING DISTRIBUTED DATA PARALLEL (DDP)



# INITIALIZE THE PROCESS

```
def setup(global_rank, world_size):  
    dist.init_process_group(backend="nccl", rank=global_rank,  
        world_size=world_size)
```

# PIN GPU TO BE USED

```
device = torch.device("cuda:" + str(local_rank))  
model = Net().to(device)
```

# ENCAPSULATE MODEL WITH DDP

```
model = nn.parallel.DistributedDataParallel(model,  
device_ids=[local_rank])
```

# SYNCHRONIZE INITIAL STATE


Handled internally by DDP across processes and nodes!



# DATA PARTITIONING

Shuffle the dataset

Partition records among  
workers



Train by sequentially reading  
the partition

After epoch is done, reshuffle  
and partition again

# DATA PARTITIONING

```
train_sampler =  
torch.utils.data.distributed.DistributedSampler(train_set,  
num_replicas=world_size, rank=global_rank)  
  
train_loader =  
torch.utils.data.DataLoader(train_set,  
batch_size=args.batch_size, sampler=train_sampler)
```

# I/O ON ONLY ON ONE WORKER

```
download = True if local_rank == 0 else False
if local_rank == 0:
    train_set = torchvision.datasets.FashionMNIST("./data",
download=download)

-----

if global_rank == 0:
    print("Epoch = {:2d}: Validation Loss = {:.5.3f},
Validation Accuracy = {:.5.3f}".format(epoch+1, v_loss,
val_accuracy[-1]))
```



**nvidia**

DEEP  
LEARNING  
INSTITUTE

[www.nvidia.com/dli](http://www.nvidia.com/dli)